

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [2]: pip install sentencepiece
```

Collecting sentencepiece

 Downloading sentencepiece-0.2.0-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (7.7 kB)

 Downloading sentencepiece-0.2.0-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (1.3 MB)

 1.3/1.3 MB 18.5 MB/s eta 0:00:00

Installing collected packages: sentencepiece

Successfully installed sentencepiece-0.2.0

Note: you may need to restart the kernel to use updated packages.

```
In [3]: pip install torch
```

```
Collecting torch
  Downloading torch-2.6.0-cp312-cp312-manylinux1_x86_64.whl.metadata (28 kB)
Collecting filelock (from torch)
  Downloading filelock-3.18.0-py3-none-any.whl.metadata (2.9 kB)
Requirement already satisfied: typing-extensions>=4.10.0 in /opt/conda/lib/python3.12/site-packages (from torch) (4.12.2)
Requirement already satisfied: networkx in /opt/conda/lib/python3.12/site-packages (from torch) (3.4.2)
Requirement already satisfied: jinja2 in /opt/conda/lib/python3.12/site-packages (from torch) (3.1.5)
Requirement already satisfied: fsspec in /opt/conda/lib/python3.12/site-packages (from torch) (2025.2.0)
Collecting nvidia-cuda-nvrtc-cu12==12.4.127 (from torch)
  Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-runtime-cu12==12.4.127 (from torch)
  Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-cupti-cu12==12.4.127 (from torch)
  Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cudnn-cu12==9.1.0.70 (from torch)
  Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cublas-cu12==12.4.5.8 (from torch)
  Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cufft-cu12==11.2.1.3 (from torch)
  Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-curand-cu12==10.3.5.147 (from torch)
  Downloading nvidia_curand_cu12-10.3.5.147-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cusolver-cu12==11.6.1.9 (from torch)
  Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cusparse-cu12==12.3.1.170 (from torch)
  Downloading nvidia_cusparse_cu12-12.3.1.170-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cusparseelt-cu12==0.6.2 (from torch)
  Downloading nvidia_cusparseelt_cu12-0.6.2-py3-none-manylinux2014_x86_64.whl.metadata (6.8 kB)
Collecting nvidia-nccl-cu12==2.21.5 (from torch)
  Downloading nvidia_nccl_cu12-2.21.5-py3-none-manylinux2014_x86_64.whl.metadata (1.8 kB)
Collecting nvidia-nvtx-cu12==12.4.127 (from torch)
  Downloading nvidia_nvtx_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.7 kB)
Collecting nvidia-nvjitlink-cu12==12.4.127 (from torch)
  Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting triton==3.2.0 (from torch)
  Downloading triton-3.2.0-cp312-cp312-manylinux_2_17_x86_64_manylinux2014_x86_64.whl.metadata (1.4 kB)
Requirement already satisfied: setuptools in /opt/conda/lib/python3.12/site-packages (from torch) (75.8.0)
Collecting sympy==1.13.1 (from torch)
  Downloading sympy-1.13.1-py3-none-any.whl.metadata (12 kB)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /opt/conda/lib/python3.12/site-packages (from sympy==1.13.1->torch) (1.3.0)
```

```
Requirement already satisfied: MarkupSafe>=2.0 in /opt/conda/lib/python3.12/site-packages (from jinja2->torch) (3.0.2)
Downloading torch-2.6.0-cp312-cp312-manylinux1_x86_64.whl (766.6 MB)
    ━━━━━━━━━━━━━━━━ 766.6/766.6 MB 26.8 MB/s eta 0:00:000
0:0100:01
Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl (363.4 MB)
    ━━━━━━━━━━━━━━ 363.4/363.4 MB 43.8 MB/s eta 0:00:000
0:0100:01
Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (1 3.8 MB)
    ━━━━━━━━━━━━ 13.8/13.8 MB 45.5 MB/s eta 0:00:000:01
0:0100:01
Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (2 4.6 MB)
    ━━━━━━━━━━ 24.6/24.6 MB 52.8 MB/s eta 0:00:000:0100:01
0:0100:01
Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (883 kB)
    ━━━━━━━━ 883.7/883.7 kB 58.9 MB/s eta 0:00:00
Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl (664.8 M B)
    ━━━━━━━━ 664.8/664.8 MB 43.0 MB/s eta 0:00:000:0100:01
0:0100:01
Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl (211.5 M B)
    ━━━━━━ 211.5/211.5 MB 68.1 MB/s eta 0:00:000:0100:01
0:0100:01
Downloading nvidia_curand_cu12-10.3.5.147-py3-none-manylinux2014_x86_64.whl (56.3 MB)
    ━━━━ 56.3/56.3 MB 48.8 MB/s eta 0:00:000:0100:01
0:0100:01
Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl (127. 9 MB)
    ━━━━ 127.9/127.9 MB 59.2 MB/s eta 0:00:000:0100:01
0:0100:01
Downloading nvidia_cusparse_cu12-12.3.1.170-py3-none-manylinux2014_x86_64.whl (20 7.5 MB)
    ━━━━ 207.5/207.5 MB 58.5 MB/s eta 0:00:000:0100:01
0:0100:01
Downloading nvidia_cusparseelt_cu12-0.6.2-py3-none-manylinux2014_x86_64.whl (150.1 MB)
    ━━━━ 150.1/150.1 MB 66.4 MB/s eta 0:00:000:0100:01
0:0100:01
Downloading nvidia_nccl_cu12-2.21.5-py3-none-manylinux2014_x86_64.whl (188.7 MB)
    ━━━━ 188.7/188.7 MB 56.5 MB/s eta 0:00:000:0100:01
0:0100:01
Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (21. 1 MB)
    ━━━━ 21.1/21.1 MB 56.5 MB/s eta 0:00:00a:0:00:01
0:0100:01
Downloading nvidia_nvtx_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (99 kB)
0:0100:01
Downloading sympy-1.13.1-py3-none-any.whl (6.2 MB)
    ━━━━ 6.2/6.2 MB 59.7 MB/s eta 0:00:00
0:0100:01
Downloading triton-3.2.0-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.w h1 (253.2 MB)
    ━━━━ 253.2/253.2 MB 56.3 MB/s eta 0:00:000:0100:01
0:0100:01
Downloading filelock-3.18.0-py3-none-any.whl (16 kB)
Installing collected packages: triton, nvidia-cusparseelt-cu12, sympy, nvidia-nvtx
```

```
-cu12, nvidia-nvjitlink-cu12, nvidia-nccl-cu12, nvidia-curand-cu12, nvidia-cufft-cu12, nvidia-cuda-runtime-cu12, nvidia-cuda-nvrtc-cu12, nvidia-cuda-cupti-cu12, nvidia-cublas-cu12, filelock, nvidia-cusparse-cu12, nvidia-cudnn-cu12, nvidia-cusolver-cu12, torch
Attempting uninstall: sympy
Found existing installation: sympy 1.13.3
Uninstalling sympy-1.13.3:
Successfully uninstalled sympy-1.13.3
Successfully installed filelock-3.18.0 nvidia-cublas-cu12-12.4.5.8 nvidia-cuda-cupti-cu12-12.4.127 nvidia-cuda-nvrtc-cu12-12.4.127 nvidia-cuda-runtime-cu12-12.4.127 nvidia-cudnn-cu12-9.1.0.70 nvidia-cufft-cu12-11.2.1.3 nvidia-curand-cu12-10.3.5.147 nvidia-cusolver-cu12-11.6.1.9 nvidia-cusparse-cu12-12.3.1.170 nvidia-cusparse-cu12-0.6.2 nvidia-nccl-cu12-2.21.5 nvidia-nvjitlink-cu12-12.4.127 nvidia-nvtx-cu12-12.4.127 sympy-1.13.1 torch-2.6.0 triton-3.2.0
Note: you may need to restart the kernel to use updated packages.
```

In [4]: pip install transformers

```

Collecting transformers
  Downloading transformers-4.50.3-py3-none-any.whl.metadata (39 kB)
Requirement already satisfied: filelock in /opt/conda/lib/python3.12/site-packages (from transformers) (3.18.0)
Collecting huggingface-hub<1.0,>=0.26.0 (from transformers)
  Downloading huggingface_hub-0.30.1-py3-none-any.whl.metadata (13 kB)
Requirement already satisfied: numpy>=1.17 in /opt/conda/lib/python3.12/site-packages (from transformers) (2.0.2)
Requirement already satisfied: packaging>=20.0 in /opt/conda/lib/python3.12/site-packages (from transformers) (24.2)
Requirement already satisfied: pyyaml>=5.1 in /opt/conda/lib/python3.12/site-packages (from transformers) (6.0.2)
Collecting regex!=2019.12.17 (from transformers)
  Downloading regex-2024.11.6-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (40 kB)
Requirement already satisfied: requests in /opt/conda/lib/python3.12/site-packages (from transformers) (2.32.3)
Collecting tokenizers<0.22,>=0.21 (from transformers)
  Downloading tokenizers-0.21.1-cp39-abi3-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (6.8 kB)
Collecting safetensors>=0.4.3 (from transformers)
  Downloading safetensors-0.5.3-cp38-abi3-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (3.8 kB)
Requirement already satisfied: tqdm>=4.27 in /opt/conda/lib/python3.12/site-packages (from transformers) (4.67.1)
Requirement already satisfied: fsspec>=2023.5.0 in /opt/conda/lib/python3.12/site-packages (from huggingface-hub<1.0,>=0.26.0->transformers) (2025.2.0)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /opt/conda/lib/python3.12/site-packages (from huggingface-hub<1.0,>=0.26.0->transformers) (4.12.2)
Requirement already satisfied: charset_normalizer<4,>=2 in /opt/conda/lib/python3.12/site-packages (from requests->transformers) (3.4.1)
Requirement already satisfied: idna<4,>=2.5 in /opt/conda/lib/python3.12/site-packages (from requests->transformers) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /opt/conda/lib/python3.12/site-packages (from requests->transformers) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in /opt/conda/lib/python3.12/site-packages (from requests->transformers) (2024.12.14)
Downloading transformers-4.50.3-py3-none-any.whl (10.2 MB)
  10.2/10.2 MB 14.6 MB/s eta 0:00:0000:
01
  Downloading huggingface_hub-0.30.1-py3-none-any.whl (481 kB)
  Downloading regex-2024.11.6-cp312-cp312-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (796 kB)
  796.9/796.9 kB 1.4 MB/s eta 0:00:00:0
0:01
  Downloading safetensors-0.5.3-cp38-abi3-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (471 kB)
  Downloading tokenizers-0.21.1-cp39-abi3-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (3.0 MB)
  3.0/3.0 MB 6.3 MB/s eta 0:00:00ta 0:0
0:01
  Installing collected packages: safetensors, regex, huggingface-hub, tokenizers, transformers
  Successfully installed huggingface-hub-0.30.1 regex-2024.11.6 safetensors-0.5.3 tokenizers-0.21.1 transformers-4.50.3
  Note: you may need to restart the kernel to use updated packages.

```

In [5]:

```

import torch
import torch.nn as nn
from torch.utils.data import DataLoader, Dataset, random_split

```

```

from transformers import MT5Tokenizer, MT5ForSequenceClassification
import pandas as pd
from tqdm import tqdm
import os

class TextDataset(Dataset):
    def __init__(self, texts, labels, tokenizer, max_length=128):
        self.texts = texts
        self.labels = labels
        self.tokenizer = tokenizer
        self.max_length = max_length

    def __len__(self):
        return len(self.texts)

    def __getitem__(self, idx):
        encoding = self.tokenizer(
            self.texts[idx],
            truncation=True,
            padding='max_length',
            max_length=self.max_length,
            return_tensors="pt"
        )
        return {
            'input_ids': encoding['input_ids'].squeeze(0),
            'attention_mask': encoding['attention_mask'].squeeze(0),
            'labels': torch.tensor(self.labels[idx], dtype=torch.long)
        }

train_df = pd.read_csv('toxic_eng/train.csv')
test_df = pd.read_csv('toxic_eng/test.csv')

train_texts = train_df['comment_text'].tolist()
train_labels = train_df['toxic'].tolist()
test_texts = test_df['comment_text'].tolist()
test_labels = test_df['toxic'].tolist()

tokenizer = MT5Tokenizer.from_pretrained("google/mt5-base")

MAX_LEN = 128
full_train_dataset = TextDataset(train_texts, train_labels, tokenizer, max_length=MAX_LEN)
test_dataset = TextDataset(test_texts, test_labels, tokenizer, max_length=MAX_LEN)

train_size = int(0.9 * len(full_train_dataset))
val_size = len(full_train_dataset) - train_size
train_dataset, val_dataset = random_split(full_train_dataset, [train_size, val_size])

train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=32)
test_loader = DataLoader(test_dataset, batch_size=32)

def save_checkpoint(model, optimizer, epoch, checkpoint_dir, model_name, best_val_loss):
    os.makedirs(checkpoint_dir, exist_ok=True)
    checkpoint_path = os.path.join(checkpoint_dir, f'{model_name}_epoch{epoch+1}')
    torch.save({
        'epoch': epoch + 1,
        'model_state_dict': model.state_dict(),
        'optimizer_state_dict': optimizer.state_dict(),
        'best_val_loss': best_val_loss
    }, checkpoint_path)

```

```

print(f"Checkpoint saved at: {checkpoint_path}")

def compute_val_loss(model, val_loader, criterion, device):
    model.eval()
    total_loss = 0
    with torch.no_grad():
        for batch in val_loader:
            input_ids = batch['input_ids'].to(device)
            attention_mask = batch['attention_mask'].to(device)
            labels = batch['labels'].to(device)
            outputs = model(input_ids=input_ids, attention_mask=attention_mask)
            loss = criterion(outputs.logits, labels)
            total_loss += loss.item()
    return total_loss / len(val_loader)

def load_checkpoint(model, optimizer, checkpoint_dir, model_name):
    checkpoint_path = os.path.join(checkpoint_dir, f"{model_name}_best.pt")
    if os.path.exists(checkpoint_path):
        checkpoint = torch.load(checkpoint_path)
        model.load_state_dict(checkpoint['model_state_dict'])
        optimizer.load_state_dict(checkpoint['optimizer_state_dict'])
        epoch = checkpoint['epoch']
        best_val_loss = checkpoint['best_val_loss']
        print(f"Resuming training from epoch {epoch + 1} with best validation loss")
        return epoch, best_val_loss
    else:
        print("No checkpoint found. Starting training from scratch.")
        return 0, float('inf')

def train_with_early_stopping(model, train_loader, val_loader, criterion, optimizer,
    best_val_loss = float('inf')
    patience_counter = 0

    epoch, best_val_loss = load_checkpoint(model, optimizer, checkpoint_dir, model_name)

    for epoch in range(epoch, num_epochs):
        model.train()
        total_loss = 0
        correct = 0
        total = 0
        loop = tqdm(train_loader, desc=f"Epoch {epoch+1}/{num_epochs}")

        for batch in loop:
            input_ids = batch['input_ids'].to(device)
            attention_mask = batch['attention_mask'].to(device)
            labels = batch['labels'].to(device)

            outputs = model(input_ids=input_ids, attention_mask=attention_mask)
            logits = outputs.logits
            loss = criterion(logits, labels)

            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

            total_loss += loss.item()
            _, predicted = torch.max(logits, dim=1)
            correct += (predicted == labels).sum().item()
            total += labels.size(0)
    
```

```

        loop.set_postfix(loss=loss.item())

    train_accuracy = correct / total
    val_loss = compute_val_loss(model, val_loader, criterion, device)
    print(f"Epoch {epoch+1}, Train Loss: {total_loss:.4f}, Train Acc: {train
        save_checkpoint(model, optimizer, epoch, checkpoint_dir, model_name, bes

    # Save best model
    if val_loss < best_val_loss:
        best_val_loss = val_loss
        patience_counter = 0
        best_path = os.path.join(checkpoint_dir, f"{model_name}_best.pt")
        torch.save({
            'epoch': epoch + 1,
            'model_state_dict': model.state_dict(),
            'optimizer_state_dict': optimizer.state_dict(),
            'best_val_loss': best_val_loss
        }, best_path)
        print(f"Best model saved at: {best_path}")
    else:
        patience_counter += 1
        print(f"Early stopping patience: {patience_counter}/{patience}")
        if patience_counter >= patience:
            print("Early stopping triggered.")
            break


def evaluate_mt5_model(model, test_loader, device):
    model.eval()
    correct = 0
    total = 0
    test_loader_tqdm = tqdm(test_loader, desc="Evaluating", leave=False)
    with torch.no_grad():
        for batch in test_loader_tqdm:
            input_ids = batch['input_ids'].to(device)
            attention_mask = batch['attention_mask'].to(device)
            labels = batch['labels'].to(device)
            outputs = model(input_ids, attention_mask)
            logits = outputs.logits
            _, preds = torch.max(logits, 1)
            correct += (preds == labels).sum().item()
            total += labels.size(0)
    accuracy = correct / total
    print(f"Test Accuracy: {accuracy:.4f}")


# Init
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = MT5ForSequenceClassification.from_pretrained("google/mt5-base", num_label
optimizer = torch.optim.Adam(model.parameters(), lr=1e-5)
criterion = nn.CrossEntropyLoss()

```

The tokenizer class you load from this checkpoint is not the same type as the class this function is called from. It may result in unexpected tokenization.
The tokenizer class you load from this checkpoint is 'T5Tokenizer'.
The class this function is called from is 'MT5Tokenizer'.
Some weights of MT5ForSequenceClassification were not initialized from the model checkpoint at google/mt5-base and are newly initialized: ['classification_head.dense.bias', 'classification_head.dense.weight', 'classification_head.out_proj.bias', 'classification_head.out_proj.weight']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
In [ ]: # Train on full dataset
train_with_early_stopping(
    model, train_loader, val_loader,
    criterion, optimizer,
    device, num_epochs=10,
    checkpoint_dir="checkpoints_full",
    model_name="mt5_toxicity_full",
    patience=3
)

# Evaluate
evaluate_mt5_model(model, test_loader, device)
```

```
In [6]: # Load best checkpoint for evaluation
def load_best_model(model_class, checkpoint_dir, model_name, device):
    checkpoint_path = os.path.join(checkpoint_dir, f"{model_name}_best.pt")
    if os.path.exists(checkpoint_path):
        print(f"Loading best model from {checkpoint_path}")
        checkpoint = torch.load(checkpoint_path, map_location=device)
        model = model_class.from_pretrained("google/mt5-base", num_labels=2)
        model.load_state_dict(checkpoint['model_state_dict'])
        model.to(device)
        return model
    else:
        raise FileNotFoundError(f"Best checkpoint not found at {checkpoint_path}")

best_model = load_best_model(MT5ForSequenceClassification, "checkpoints_full", "evaluate_mt5_model(best_model, test_loader, device))
```

Loading best model from checkpoints_full/mt5_toxicity_full_best.pt

Some weights of MT5ForSequenceClassification were not initialized from the model checkpoint at google/mt5-base and are newly initialized: ['classification_head.dense.bias', 'classification_head.dense.weight', 'classification_head.out_proj.bias', 'classification_head.out_proj.weight']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Test Accuracy: 0.9405

```
In [8]: # Load best checkpoint for evaluation
def load_best_model(model_class, checkpoint_dir, model_name, device):
    checkpoint_path = os.path.join(checkpoint_dir, f"{model_name}.pt")
    if os.path.exists(checkpoint_path):
        print(f"Loading best model from {checkpoint_path}")
        checkpoint = torch.load(checkpoint_path, map_location=device)
        model = model_class.from_pretrained("google/mt5-base", num_labels=2)
        model.load_state_dict(checkpoint['model_state_dict'])
        model.to(device)
```

```

        return model
    else:
        raise FileNotFoundError(f"Best checkpoint not found at {checkpoint_path}")

# Load best model and evaluate it on test set
best_model = load_best_model(MT5ForSequenceClassification, "checkpoints_full", "evaluate_mt5_model(best_model, test_loader, device)

```

Loading best model from checkpoints_full/mt5_toxicity_full_epoch10.pt

Some weights of MT5ForSequenceClassification were not initialized from the model checkpoint at google/mt5-base and are newly initialized: ['classification_head.dense.bias', 'classification_head.dense.weight', 'classification_head.out_proj.bias', 'classification_head.out_proj.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Test Accuracy: 0.9423

```
In [ ]: import os

checkpoint_dir = "checkpoints"
files = os.listdir(checkpoint_dir)
print("Obsah priečinka checkpoints:")
for file in files:
    print(f"- {file}")
```

```
In [9]: import torch
from transformers import MT5ForSequenceClassification
from sklearn.metrics import classification_report, confusion_matrix
from tqdm import tqdm
import os
# Load full dataset
train_df = pd.read_csv('toxic_eng/train.csv')
test_df = pd.read_csv('toxic_eng/test.csv')

train_texts = train_df['comment_text'].tolist()
train_labels = train_df['toxic'].tolist()
test_texts = test_df['comment_text'].tolist()
test_labels = test_df['toxic'].tolist()

# Tokenizer
tokenizer = MT5Tokenizer.from_pretrained("google/mt5-base")
# Načítanie checkpointu
def load_best_model(model_class, checkpoint_dir, model_name, device):
    checkpoint_path = os.path.join(checkpoint_dir, f"{model_name}.pt")
    if os.path.exists(checkpoint_path):
        print(f"Loading best model from {checkpoint_path}")
        checkpoint = torch.load(checkpoint_path, map_location=device)
        model = model_class.from_pretrained("google/mt5-base", num_labels=2)
        model.load_state_dict(checkpoint['model_state_dict'])
        model.to(device)
        return model
    else:
        raise FileNotFoundError(f"Best checkpoint not found at {checkpoint_path}")

# Vyhodnotenie
def evaluate_mt5_model(model, test_loader, device):
    model.eval()
    all_preds = []
```

```

all_labels = []

with torch.no_grad():
    for batch in tqdm(test_loader, desc="Vyhodnocovanie"):
        input_ids = batch['input_ids'].to(device)
        attention_mask = batch['attention_mask'].to(device)
        labels = batch['labels'].to(device)

        outputs = model(input_ids=input_ids, attention_mask=attention_mask)
        logits = outputs.logits
        preds = torch.argmax(logits, dim=1)

        all_preds.extend(preds.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

print(f"\nPresnosť na testovacích dátach: {torch.sum(torch.tensor(all_preds) == all_labels).item() / len(all_labels)}")
print("\nKlasifikačná správa:")
print(classification_report(all_labels, all_preds, digits=4))
print("\nConfusion Matrix:")
print(confusion_matrix(all_labels, all_preds))

# === Inicializácia ===
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
best_model = load_best_model(MT5ForSequenceClassification, "checkpoints_full", "checkpoints_full")

# === Vyhodnotenie ===
evaluate_mt5_model(best_model, test_loader, device)

```

The tokenizer class you load from this checkpoint is not the same type as the class this function is called from. It may result in unexpected tokenization.
The tokenizer class you load from this checkpoint is 'T5Tokenizer'.
The class this function is called from is 'MT5Tokenizer'.

Loading best model from checkpoints_full/mt5_toxicity_full_epoch10.pt

Some weights of MT5ForSequenceClassification were not initialized from the model checkpoint at google/mt5-base and are newly initialized: ['classification_head.dense.bias', 'classification_head.dense.weight', 'classification_head.out_proj.bias', 'classification_head.out_proj.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Vyhodnocovanie: 100%|██████████| 625/625 [17:05<00:00, 1.64s/it]

Presnosť na testovacích dátach: 0.9423

Klasifikačná správa:

	precision	recall	f1-score	support
0	0.9623	0.9206	0.9410	10000
1	0.9239	0.9639	0.9435	10000
accuracy			0.9423	20000
macro avg	0.9431	0.9423	0.9422	20000
weighted avg	0.9431	0.9423	0.9422	20000

Confusion Matrix:

[[9206 794]
[361 9639]]

In []:

